

Harris Corner Detector

Team #8

陳昱銓 陳冠豪 楊煥佑

https://github.com/yqcheneee/ACA_21S_final

Outline

- Introduction
- System & Algorithms
- Optimization & Experimental Result
- Conclusion

Outline

- Introduction
- System & Algorithms
- Optimization & Experimental Result
- Conclusion

Introduction

- Feature extraction is a common technique in computer vision and machine learning application.
- Harris corner detector is a method to extract features of corners in the image.
- In this project, we exploit parallel computation to accelerate the process of harris corner detector

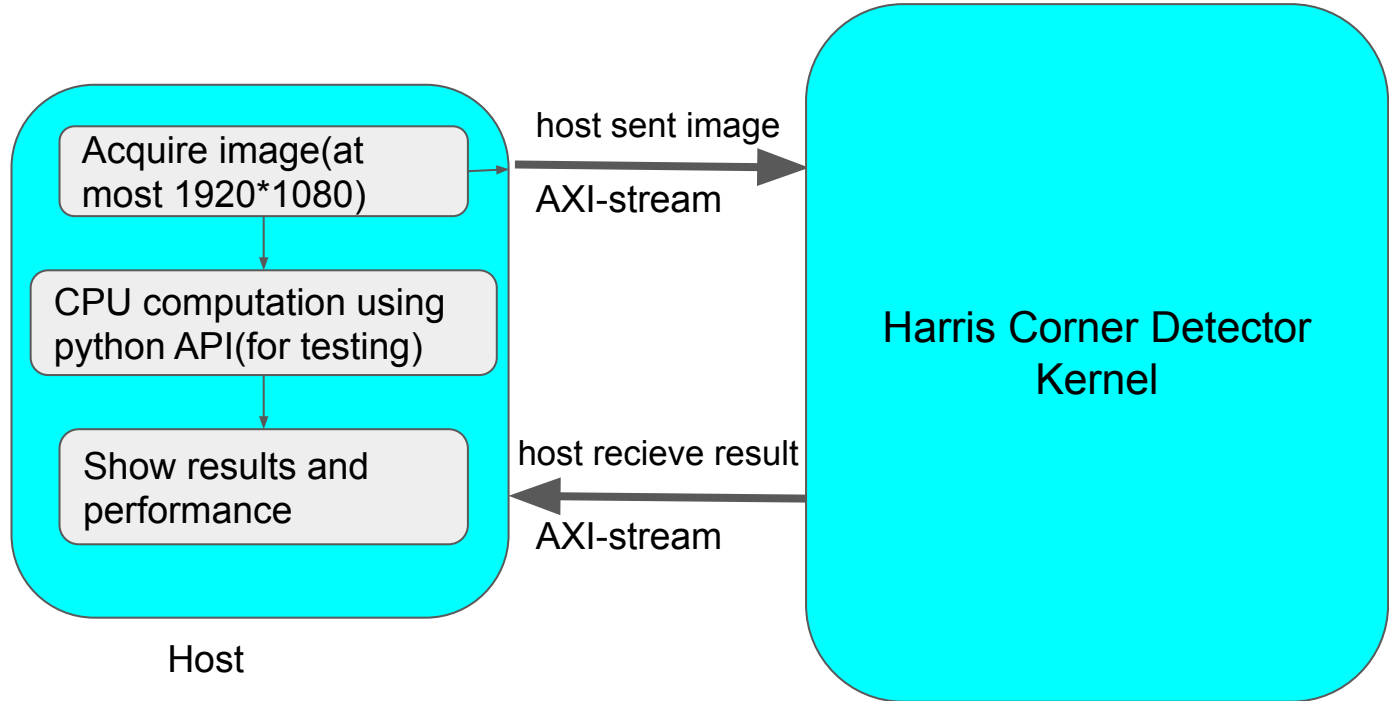
Introduction

- We use Xilinx Vivado HLS to implement hardware design.
- Hardware can be deployed on PYNQ-Z2 board.
- We explore some optimization methods.
- Our target specification: performance(latency) and precision
 - We try to keep calculation error low enough and optimize performance as much as possible

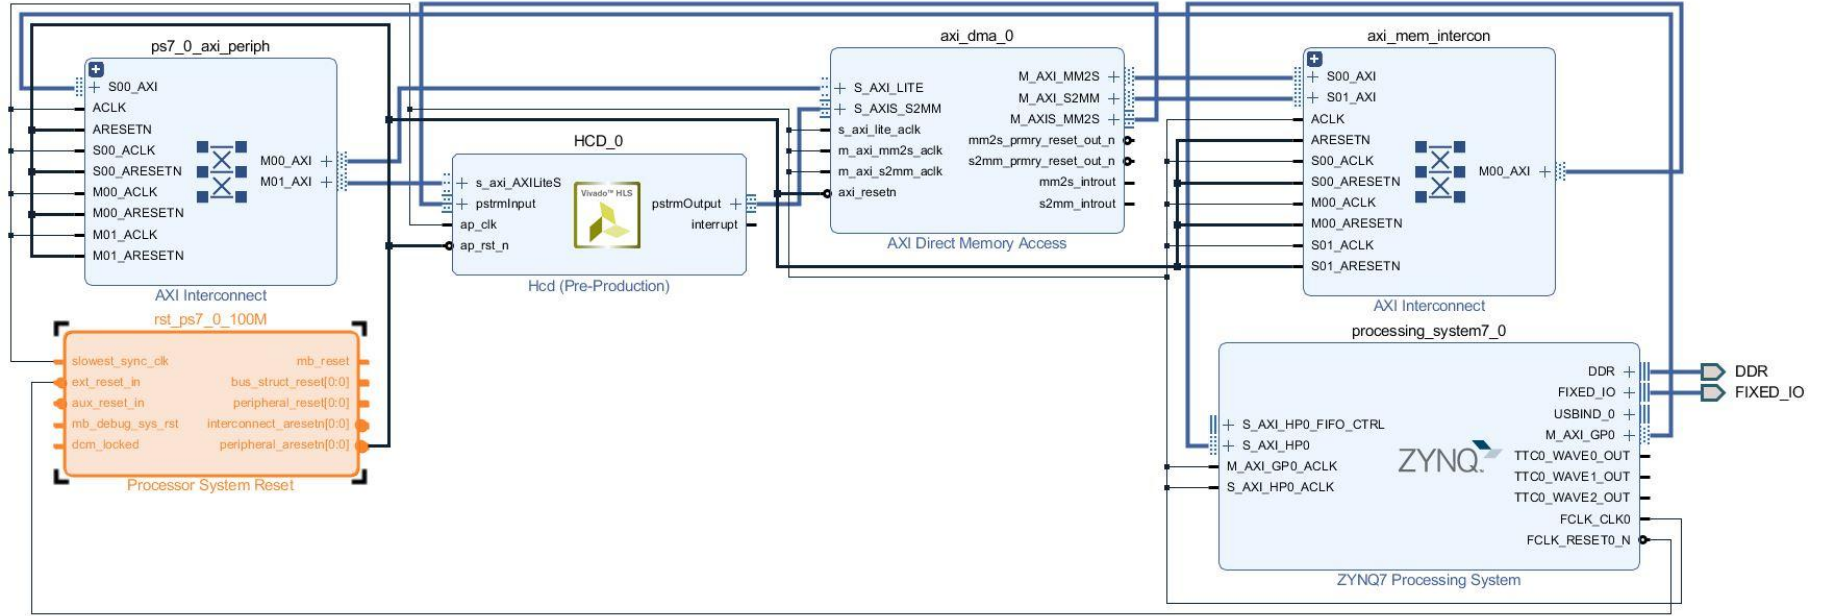
Outline

- Introduction
- System & Algorithms
- Optimization & Experimental Result
- Conclusion

System Overview




PL Block Diagram



Harris Corner Detector Algorithm

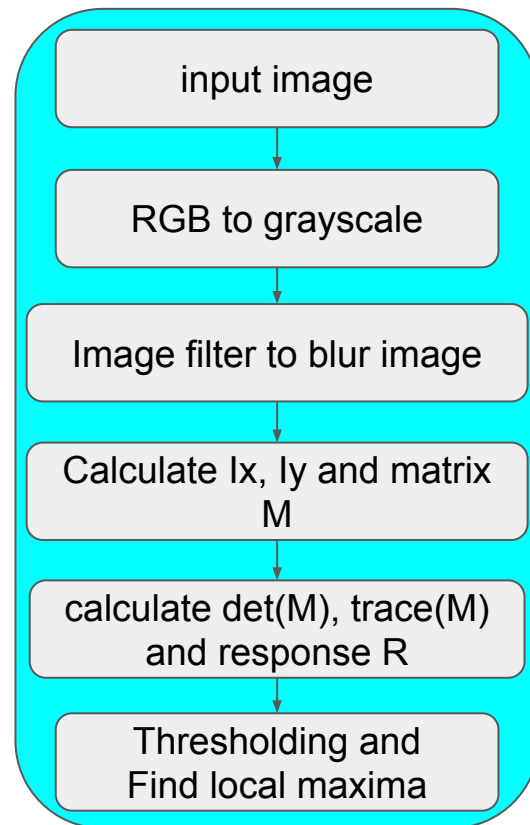
- Image filter
 - gaussianBlur
- I_x, I_y
 - gradients on x, y direction

$$M = \sum_{(x,y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$





$$R = \det(M) - k(\text{trace}(M))^2 = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2$$

Harris Corner Detector Kernel



Harris Corner Detector Example



original image $\xrightarrow{\text{image filter \& matrix operation}}$ response

thresholding

local maxima
(final result)

Outline


- Introduction
- System & Algorithms
- Optimization & Experimental Result
- Conclusion

Optimize top function (HCD kernel)

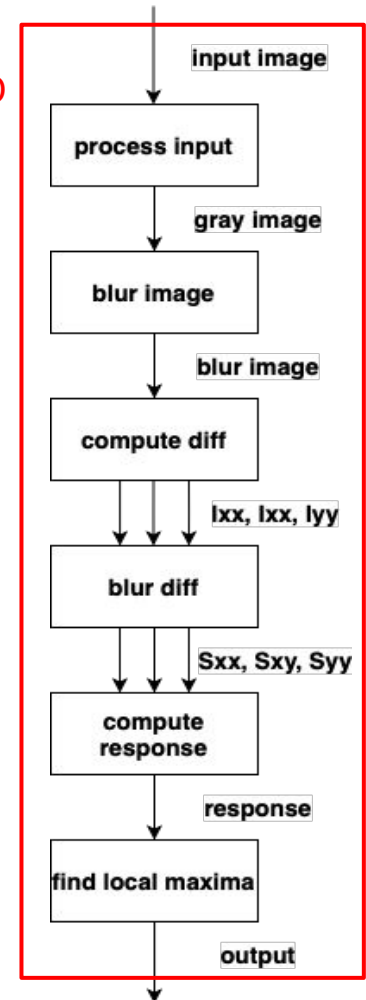
- HLS pragma Dataflow
 - Make the functions operate in parallel to achieve high throughput

```
void HCD(stream_t* pstrmInput, stream_t* pstrmOutput,  
         reg32_t row, reg32_t col)  
{  
    #pragma HLS DATAFLOW  
    process_input(pstrmInput, &stream_gray, row, col);  
  
    blur_img(&stream_gray, &stream_blur, row, col);  
  
    compute_dif(&stream_blur, &stream_Ixx, &stream_Iyy,  
               &stream_Ixy, row, col);  
  
    blur_diff(&stream_Ixx, &stream_Iyy, &stream_Ixy,  
              &stream_Sxx, &stream_Syy, &stream_Sxy, row, col);  
  
    compute_det_trace(&stream_Sxx, &stream_Syy, &stream_Sxy,  
                     &stream_response, row, col);  
  
    find_local_maxima(&stream_response, pstrmOutput, row, col);  
}
```

```
void sub_function(stream* input, stream* output,  
                 int row, int col)  
{  
    for (i in row) {  
        for (j in col) {  
            tmp = input-> read()  
  
            if (condition) {  
                operation on tmp  
            }  
  
            output-> write(tmp);  
        }  
    }  
}
```



HCD



Optimization inside sub-functions

- Reduce resources usage
 - Merge all loops into one loop to reduce the number of unexpectable states in finite-state machine

```
template<typename T>
void blur_img(hls::stream<T>& stream_gray, hls::stream<T>& stream_out)
{
    // initial read
    for (i = 0 ; i < 2 ; i++) {
        for (j = 0 ; j < MAX_WIDTH ; j++) { ...
        }

        // i = 0
        for (j = 1; j < MAX_WIDTH; ++j) { ...

        // i = 1 ~ MAX_HEIGHT-2
        for (i = 1; i < MAX_HEIGHT - 1; ++i) {
            for (j = 1; j < MAX_WIDTH; ++j) { ...
        }

        // i = MAX_HEIGHT
        for (j = 1; j < MAX_WIDTH; ++j) { ...
    }
}
```

* Multiplexer:

Name	LUT
ColIndex_assign_5_reg_6519	9
ColIndex_assign_6_reg_6316	9
ColIndex_assign_reg_6159	9
ap_NS_fsm	1761
buf_M_0_V_address0	597
buf_M_0_V_address1	597
buf_M_0_V_d0	585
buf_M_0_V_d1	581
buf_M_1_V_address0	1185
buf_M_1_V_address1	1189
buf_M_1_V_d0	589
buf_M_1_V_d1	545
buf_M_2_V_address0	585
buf_M_2_V_address1	585

	LUT
	9
	9
	9
	1761
	597
	597
	585
	581
	1185
	1189
	589
	545
	585
	585

Optimization inside sub-functions

- Using **ap_fixed** instead of floating point
 - Improve both utilization and timing
 - float operations consume a lot of time and resources
 - Inside functions
 - blur image (gaussian blur)
 - compute response

	latency (cycle)	FF	LUT
float	159	960	1803
ap_fixed	33	60	245

```
template<typename P, typename W>
P Gaussian_filter_1(W* window)
{
    float sum = 0;
    P pixel = 0;

    float op[3][3] =
    { ...

    for (i = 0; i < 3; i++) {
        for (j = 0; j < 3; j++) {
            sum += window->getval(i,j) * op[i][j];
        }
    }
    pixel = P(sum);
    return pixel;
}
```



```
template<typename P, typename W>
P Gaussian_filter_1(W* window)
{
    ap_fixed<22, 16> sum = 0;
    P pixel = 0;

    ap_fixed<8, 2> op[3][3] =
    { ...

    for (i = 0; i < 3; i++) {
        for (j = 0; j < 3; j++) {
            sum += window->getval(i,j) * op[i][j];
        }
    }
    pixel = P(sum);
    return pixel;
}
```

Array Partition- Gaussian Blur (Compute weighted sum of a window)

- Complete partition window and const array
- Entire loop can be unrolled
- Expand entire loop
- Reduce one cycle compared to unrolled loop

```
ap_fixed<12, 1> op[3][3] =  
{  
    {0.0751136, 0.123841, 0.0751136},  
    {0.123841, 0.20418, 0.123841},  
    {0.0751136, 0.123841, 0.0751136}  
};  
#pragma HLS array_partition variable=op complete  
  
for (i = 0; i < 3; i++) {  
    #pragma HLS unroll  
    for (j = 0; j < 3; j++) {  
        #pragma HLS unroll  
        sum += window->getval(i,j) * op[i][j];  
    }  
}
```

```
ap_fixed<12, 1> op[3][3] =  
{  
    {0.0751136, 0.123841, 0.0751136},  
    {0.123841, 0.20418, 0.123841},  
    {0.0751136, 0.123841, 0.0751136}  
};  
#pragma HLS array_partition variable=op complete  
#pragma HLS expression_balance  
  
sum = window->getval(0,0) * op[0][0] +  
    window->getval(0,1) * op[0][1] +  
    window->getval(0,2) * op[0][2] +  
    window->getval(1,0) * op[1][0] +  
    window->getval(1,1) * op[1][1] +  
    window->getval(1,2) * op[1][2] +  
    window->getval(2,0) * op[2][0] +  
    window->getval(2,1) * op[2][1] +  
    window->getval(2,2) * op[2][2];
```

Loop Unrolling - local maxima (Find local maximum on a 5*5 region)

- Consider border condition
- Original version that can't be unrolled
- Revised version that can be unrolled

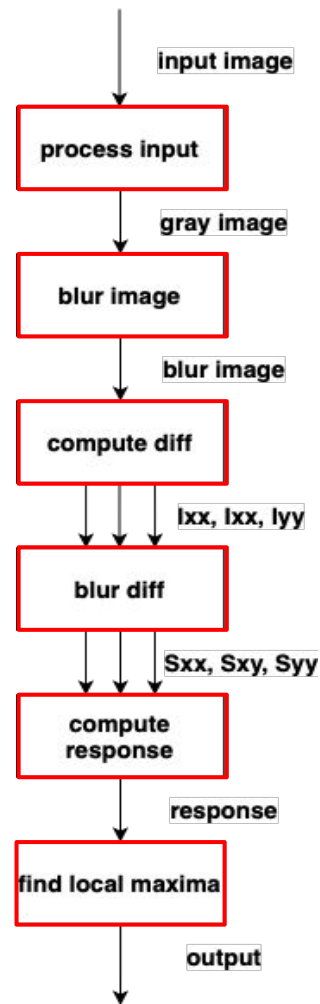
```
if (center_pixel != 0) {
    input.data = 1;
    d_bound = (i-4 < 0) ? i : 4;
    l_bound = (j-4 < 0) ? 0 : j-4;
    r_bound = (j >= col) ? col-1 : j;
    for(si = 0 ; si <= d_bound; si++) {
        for(sj = l_bound; sj <= r_bound; sj++) {
            if(response_buf.getval(si, sj) > center_pixel)
                input.data = 0;
                break;
        }
    }
}
```

```
if (center_pixel != 0) {
    input.data = 1;
    for(sj = 0; sj < 5; sj++) {
        #pragma HLS unroll
        for(si = 0 ; si < 5; si++) {
            #pragma HLS unroll
            // handle bound at getval function
            if(response_buf.getval(si, j-sj) > center_pixel) {
                input.data = 0;
                break;
            }
        }
    }
}
```


Optimization inside sub-functions

- Improve timing in all sub-functions
 - all functions have the following structures
 - use pipeline to enhance the throughput

```
void sub_function(stream* input, stream* output,  
                  int row, int col)  
{  
    for (i in row) {  
        for (j in col) {  
            # pragma HLS pipeline enable_flush  
            tmp = input-> read()  
  
            if (condition) {  
                operation on tmp  
            }  
  
            output-> write(tmp);  
        }  
    }  
}
```


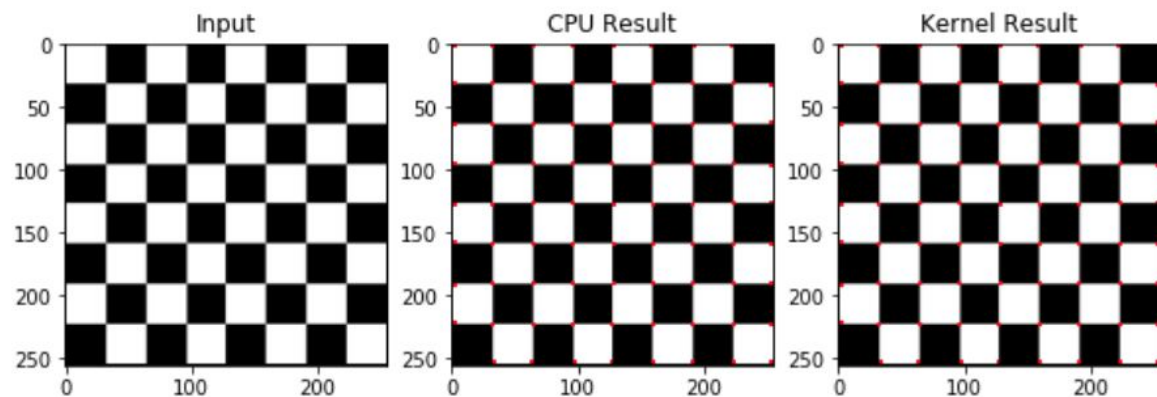


	latency			utilization			
	Synthesis report	Run time 1 (256*256)	Run time 2 (1920*1080)	BRAM_18k	DSP48 E	FF	LUT
no pragma	58 ms	46.96 ms [1]	1472 ms [1]	120	26	4433	11199
unroll	5.290 ms	5.860 ms [1]	171.5 ms [1]	80	92	9299	16640
pipeline + unroll	2.663 ms	3.316 ms [1]	83.51 ms [1]	80	127	16725	18057
CPU runtime		235 ms [2]	2425 ms [2]				

[1] run under pynq

[2] run with python under Intel(R) Xeon(R) CPU E5-2650 v2 @ 2.60GHz

Experimental results



Outline

- Introduction
- System & Algorithms
- Optimization & Experimental Result
- Conclusion

Conclusion

- We use Xilinx tools to implement Harris Corner Detector into hardware design.
- Optimize the hls code to meet the resource constraint and shorten the latency.
- Deploy our design on PYNQ-Z2 board.
- Kernel function speed up over 30x compare to CPU.